

Milestone_1_2_3_FinalSubmisison

August 12, 2023

Milestone 1

The housing market has always been in a constant state of change and hard to predict. Using housing data from Zillow, an analysis on sales over the last 10 years can be conducted to see if there can be any identification of states that are increasing in price. More importantly, the analysis can include the trends over time. The problem this analysis will solve is to research the trends of sales throughout time in different cities and states, to see where properties are the highest valued over time and how time has affected the sale prices.

```
[1]: import pandas as pd
```

```
[2]: # Read data
data = pd.read_csv('Sale_Prices_City.csv')
```

Here we can see a sample of the data.

```
[3]: data.head()
```

```
[3]:
```

	Unnamed: 0	RegionID	RegionName	StateName	SizeRank	2008-03	\
0	0	6181	New York	New York	1	NaN	
1	1	12447	Los Angeles	California	2	507600.0	
2	2	39051	Houston	Texas	3	138400.0	
3	3	17426	Chicago	Illinois	4	325100.0	
4	4	6915	San Antonio	Texas	5	130900.0	

	2008-04	2008-05	2008-06	2008-07	...	2019-06	2019-07	2019-08	\
0	NaN	NaN	NaN	NaN	...	563200.0	570500.0	572800.0	
1	489600.0	463000.0	453100.0	438100.0	...	706800.0	711800.0	717300.0	
2	135500.0	132200.0	131000.0	133400.0	...	209700.0	207400.0	207600.0	
3	314800.0	286900.0	274600.0	268500.0	...	271500.0	266500.0	264900.0	
4	131300.0	131200.0	131500.0	131600.0	...	197100.0	198700.0	200200.0	

	2019-09	2019-10	2019-11	2019-12	2020-01	2020-02	2020-03
0	569900.0	560800.0	571500.0	575100.0	571700.0	568300.0	573600.0
1	714100.0	711900.0	718400.0	727100.0	738200.0	760200.0	NaN
2	207000.0	211400.0	211500.0	217700.0	219200.0	223800.0	NaN
3	265000.0	264100.0	264300.0	270000.0	281400.0	302900.0	309200.0
4	200800.0	203400.0	203800.0	205400.0	205400.0	208300.0	NaN

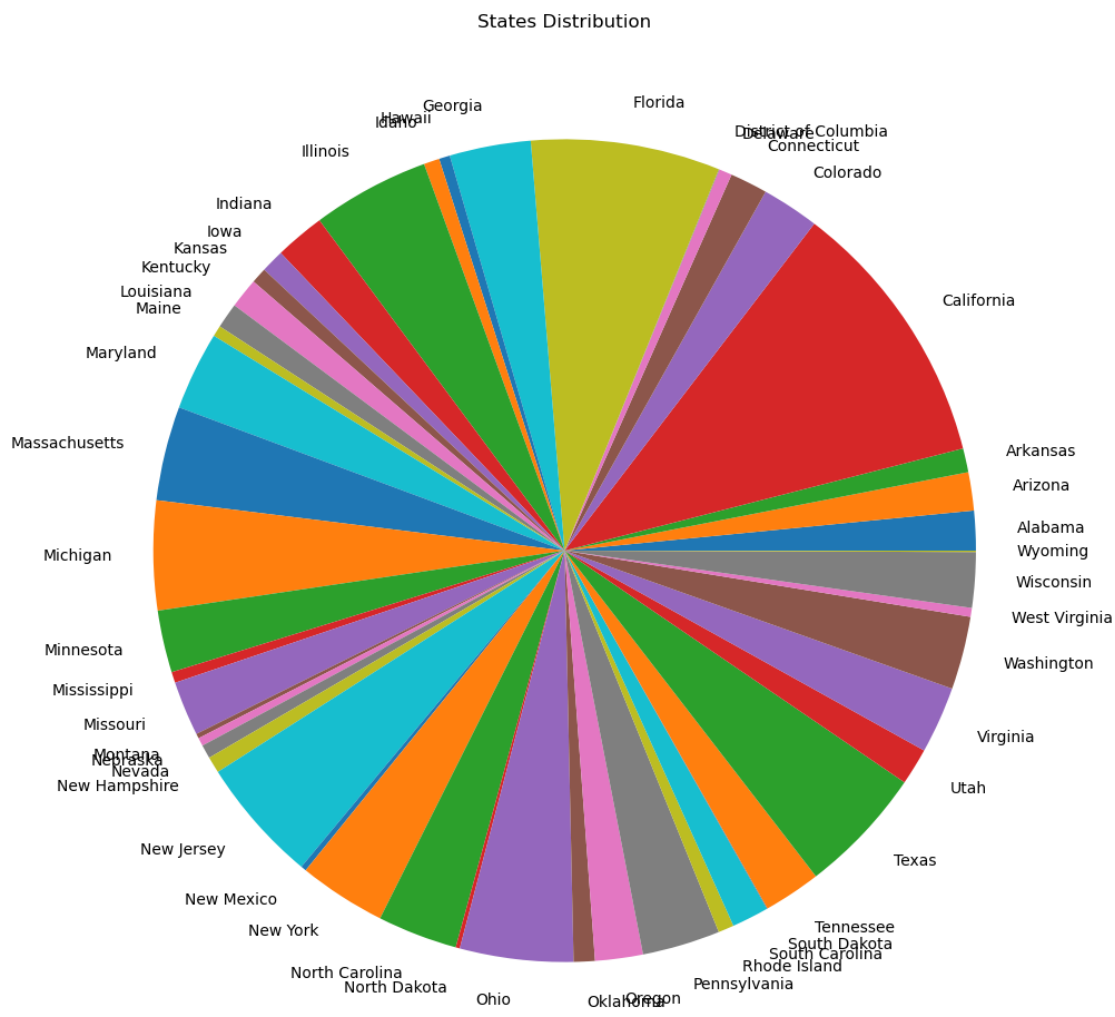
[5 rows x 150 columns]

```
[4]: # Import libraries
import matplotlib.pyplot as plt
import matplotlib
```

The first graph shown will be a pie chart. This chart can give an overview of the distribution of States to see which have the most density (sales).

```
[5]: # Create data only containing StateName
pie_data = data.groupby('StateName').size()

# Make the pie plot with pandas
pie_data.plot(kind='pie', subplots=True, figsize=(12,12), labeldistance=1.1)
plt.title("States Distribution")
plt.ylabel("")
plt.show()
```

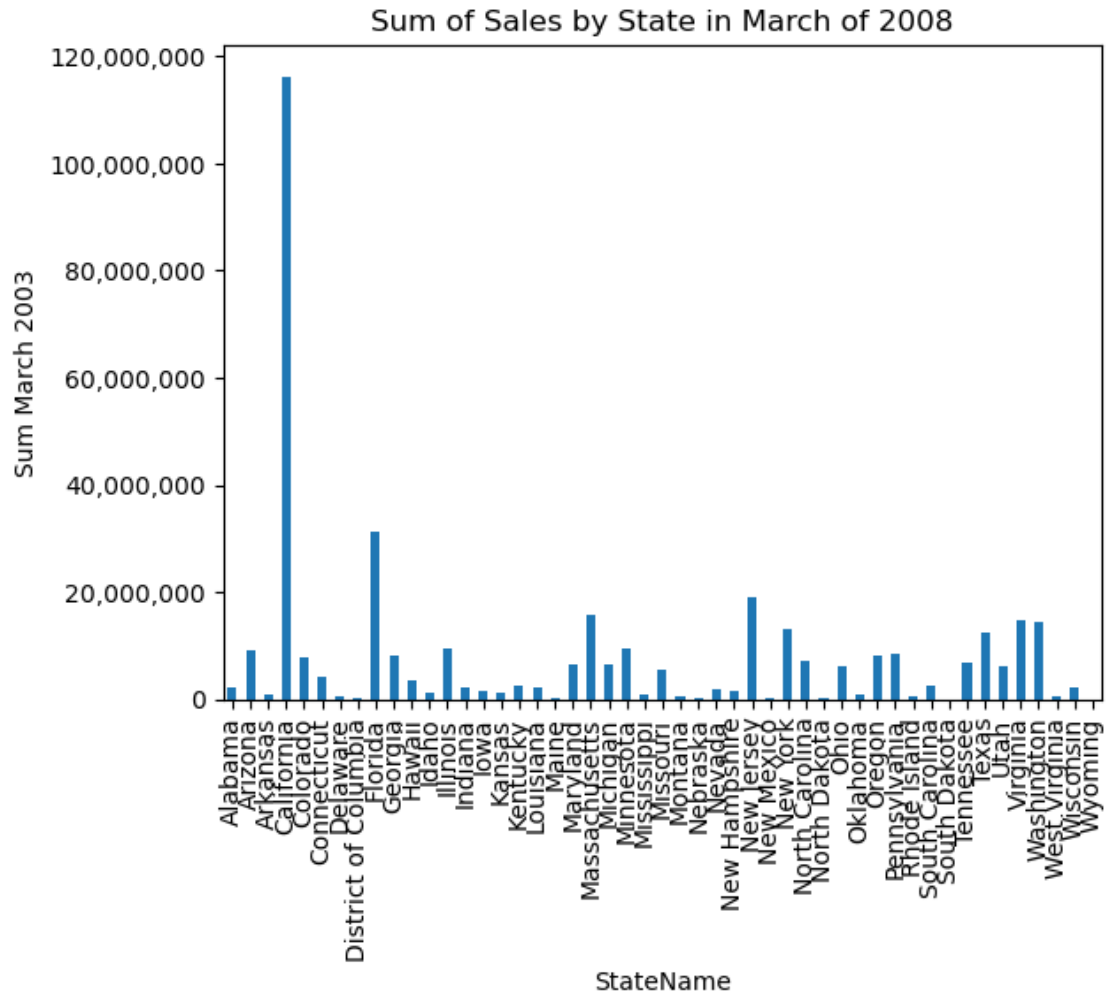


When analyzing this chart, we can see that there are a few states that stick out in terms of size. These are California, Florida, and Massachusetts.

The next graph that will be created, will be the sales of the first data point, March of 2008. This can be a good starting point to being to understand the changes of price over time.

```
[6]: # Groups by state and sums the prices of March 2008
df_grouped = data.groupby('StateName').sum()['2008-03']

[7]: # Creates plot using grouped data
ax = df_grouped.plot(kind='bar')
# Formats numbers to regular notation instead of scientific
ax.get_yaxis().set_major_formatter(
    matplotlib.ticker.FuncFormatter(lambda x, p: format(int(x), ',')))
# Creates Labels
plt.xlabel('StateName')
plt.ylabel('Sum March 2003')
plt.title('Sum of Sales by State in March of 2008')
# Rotates X labels 90 degrees
plt.xticks(rotation=90)
# Displays plot
plt.show()
```



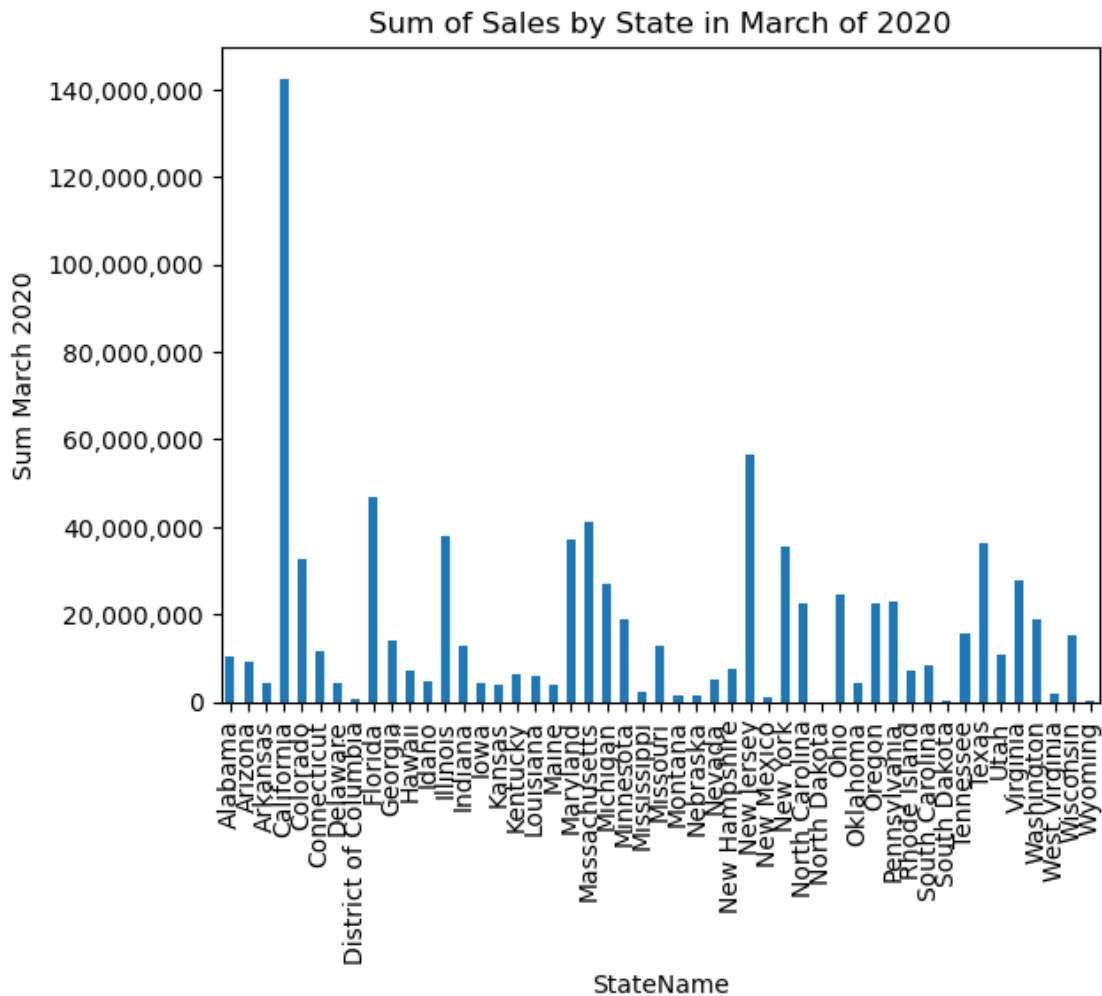
This graph confirms some of the assumptions seen in the previous plot. This shows that in the month of March in 2008, the sum of total sales are the highest in California, Florida, New Jersey, and Massachusetts.

Next, a graph using more recent sales, March of 2020.

```
[8]: # Grouping by State and sum of March of 2020
df_grouped2 = data.groupby('StateName').sum()['2020-03']
```

```
[9]: # Creates plot using the grouped data
ax = df_grouped2.plot(kind='bar')
# Formats numbers to regular notation instead of scientific
ax.get_yaxis().set_major_formatter(
    matplotlib.ticker.FuncFormatter(lambda x, p: format(int(x), ',')))
# Creates Labels
plt.xlabel('StateName')
```

```
plt.ylabel('Sum March 2020')
plt.title('Sum of Sales by State in March of 2020')
# Rotates X-Labels 90 degrees
plt.xticks(rotation=90)
# Displays Plot
plt.show()
```



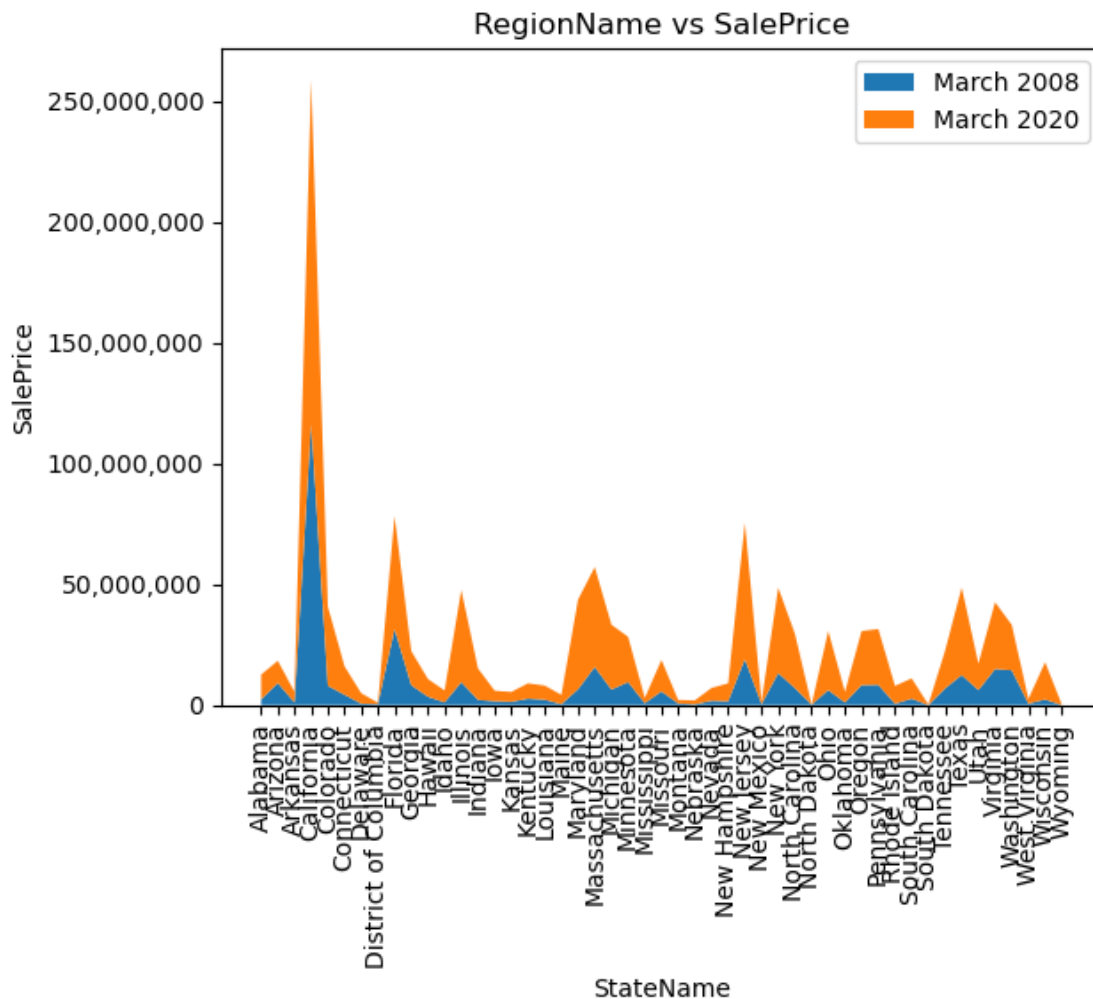
As expected, we see a rise in almost all states that are reporting data. Now, there are more states that share a similar amount in sales. Next, a stack plot can be created to see the two data sets overlaid.

```
[10]: # Create figure and axes
fig, ax = plt.subplots()
# Plot data
ax.stackplot(df_grouped.index, df_grouped.values, df_grouped2.values,
             labels=['March 2008', 'March 2020'])
```

```

# Formats numbers to regular notation instead of scientific
ax.get_yaxis().set_major_formatter(
    matplotlib.ticker.FuncFormatter(lambda x, p: format(int(x), ',')))
# Set labels
ax.set_xlabel('StateName')
ax.set_ylabel('SalePrice')
ax.set_title('RegionName vs SalePrice')
# Rotates X-Labels 90 Degrees
plt.xticks(rotation=90)
# Creates Legend
ax.legend()
# Displays plot
plt.show()

```



The graph above shows the changes in sales side by side from March 2008 to March 2020.

In more recent times, it can be seen that sales are higher than they were in 2008. The graphs created have demonstrated those changes and show that certain states increased more than others. With this information, continuing analysis can be done to see which states have increased the most over time. These graphs begin to give some insight but more research will be conducted to see where the highest valued areas are located and where the increase is the highest.

Milestone 2

The dataset used for modeling will be cleaned and new features will be created to aid with the model.

```
[11]: import pandas as pd
      # Formats numbers to regular view instead of scientific notation
      pd.set_option('display.float_format', lambda x: '%.3f' % x)
```

```
[12]: # Read data
      data2 = pd.read_csv('Sale_Prices_City.csv')
```

```
[13]: data2.head()
```

```
[13]: Unnamed: 0  RegionID  RegionName  StateName  SizeRank  2008-03  \
0            0      6181    New York    New York         1      NaN
1            1     12447  Los Angeles  California         2  507600.000
2            2     39051    Houston     Texas          3  138400.000
3            3     17426    Chicago    Illinois         4  325100.000
4            4      6915  San Antonio     Texas          5  130900.000

      2008-04  2008-05  2008-06  2008-07  ...  2019-06  2019-07  \
0          NaN        NaN        NaN        NaN  ...  563200.000  570500.000
1  489600.000  463000.000  453100.000  438100.000  ...  706800.000  711800.000
2  135500.000  132200.000  131000.000  133400.000  ...  209700.000  207400.000
3  314800.000  286900.000  274600.000  268500.000  ...  271500.000  266500.000
4  131300.000  131200.000  131500.000  131600.000  ...  197100.000  198700.000

      2019-08  2019-09  2019-10  2019-11  2019-12  2020-01  \
0  572800.000  569900.000  560800.000  571500.000  575100.000  571700.000
1  717300.000  714100.000  711900.000  718400.000  727100.000  738200.000
2  207600.000  207000.000  211400.000  211500.000  217700.000  219200.000
3  264900.000  265000.000  264100.000  264300.000  270000.000  281400.000
4  200200.000  200800.000  203400.000  203800.000  205400.000  205400.000

      2020-02  2020-03
0  568300.000  573600.000
1  760200.000        NaN
2  223800.000        NaN
3  302900.000  309200.000
4  208300.000        NaN
```

```
[5 rows x 150 columns]
```

The initial dataset has three columns that provide information that will not be needed for this study. These are the columns “Unnamed:”, “RegionId” and “SizeRank”.

```
[14]: data2 = data2.drop(columns='Unnamed: 0')
```

```
[15]: data2 = data2.drop(columns='RegionID')
```

```
[16]: data2 = data2.drop(columns='SizeRank')
```

```
[17]: data2.head()
```

```
[17]:
```

	RegionName	StateName	2008-03	2008-04	2008-05	2008-06	\
0	New York	New York	NaN	NaN	NaN	NaN	
1	Los Angeles	California	507600.000	489600.000	463000.000	453100.000	
2	Houston	Texas	138400.000	135500.000	132200.000	131000.000	
3	Chicago	Illinois	325100.000	314800.000	286900.000	274600.000	
4	San Antonio	Texas	130900.000	131300.000	131200.000	131500.000	

	2008-07	2008-08	2008-09	2008-10	...	2019-06	2019-07	\
0	NaN	NaN	NaN	NaN	...	563200.000	570500.000	
1	438100.000	423200.000	407800.000	396300.000	...	706800.000	711800.000	
2	133400.000	135400.000	138000.000	136400.000	...	209700.000	207400.000	
3	268500.000	264400.000	267100.000	268400.000	...	271500.000	266500.000	
4	131600.000	132300.000	131600.000	131800.000	...	197100.000	198700.000	

	2019-08	2019-09	2019-10	2019-11	2019-12	2020-01	\
0	572800.000	569900.000	560800.000	571500.000	575100.000	571700.000	
1	717300.000	714100.000	711900.000	718400.000	727100.000	738200.000	
2	207600.000	207000.000	211400.000	211500.000	217700.000	219200.000	
3	264900.000	265000.000	264100.000	264300.000	270000.000	281400.000	
4	200200.000	200800.000	203400.000	203800.000	205400.000	205400.000	

	2020-02	2020-03
0	568300.000	573600.000
1	760200.000	NaN
2	223800.000	NaN
3	302900.000	309200.000
4	208300.000	NaN

[5 rows x 147 columns]

To help with identification, the column ‘RegionName’ can be renamed to CityName.

```
[18]: data2 = data2.rename(columns={'RegionName': 'CityName'})
```

```
[19]: data2.head()
```



```
[19]:
```

	CityName	StateName	2008-03	2008-04	2008-05	2008-06	\
0	New York	New York	NaN	NaN	NaN	NaN	
1	Los Angeles	California	507600.000	489600.000	463000.000	453100.000	
2	Houston	Texas	138400.000	135500.000	132200.000	131000.000	
3	Chicago	Illinois	325100.000	314800.000	286900.000	274600.000	
4	San Antonio	Texas	130900.000	131300.000	131200.000	131500.000	

	2008-07	2008-08	2008-09	2008-10	...	2019-06	2019-07	\
0	NaN	NaN	NaN	NaN	...	563200.000	570500.000	
1	438100.000	423200.000	407800.000	396300.000	...	706800.000	711800.000	
2	133400.000	135400.000	138000.000	136400.000	...	209700.000	207400.000	
3	268500.000	264400.000	267100.000	268400.000	...	271500.000	266500.000	
4	131600.000	132300.000	131600.000	131800.000	...	197100.000	198700.000	

	2019-08	2019-09	2019-10	2019-11	2019-12	2020-01	\
0	572800.000	569900.000	560800.000	571500.000	575100.000	571700.000	
1	717300.000	714100.000	711900.000	718400.000	727100.000	738200.000	
2	207600.000	207000.000	211400.000	211500.000	217700.000	219200.000	
3	264900.000	265000.000	264100.000	264300.000	270000.000	281400.000	
4	200200.000	200800.000	203400.000	203800.000	205400.000	205400.000	

	2020-02	2020-03
0	568300.000	573600.000
1	760200.000	NaN
2	223800.000	NaN
3	302900.000	309200.000
4	208300.000	NaN

[5 rows x 147 columns]

The columns provide the data by month which can be used for analysis but analyzing year over year can give a quicker view. The columns need to be split by year and then all the data needs to be added by each similar year. The function below splits the column name and then adds the values in the dataframe.

```
[20]: # Function that splits column name and then adds values
def add_columns_by_year(df):
    # Split the year from the column
    years = [col.split('-')[0] for col in df.columns]
    # Group the columns by year and sum them
    grouped = df.groupby(years, axis=1).sum()
    return grouped
```

```
[21]: # Call to function
result = add_columns_by_year(data2)
result.head()
```

```
[21]:
```

	2008	2009	2010	2011	2012	2013	\
0	0.000	0.000	457300.000	5597200.000	5697000.000	5804300.000	
1	4333600.000	4093300.000	4244700.000	4079000.000	4165700.000	5326100.000	
2	1339800.000	1602700.000	1615800.000	1615800.000	1679400.000	1897400.000	
3	2782900.000	2741900.000	2462400.000	2192000.000	2178800.000	2536200.000	
4	1309900.000	1588500.000	1576700.000	1579600.000	1628300.000	1738900.000	

	2014	2015	2016	2017	2018	2019	\
0	5870900.000	6219500.000	6444300.000	6510700.000	6665700.000	6793500.000	
1	5515300.000	5917500.000	6367300.000	6817700.000	7667800.000	8478200.000	
2	2025300.000	2051800.000	2171900.000	2293200.000	2367300.000	2507100.000	
3	2867400.000	2840100.000	2631500.000	2722900.000	2960600.000	3234000.000	
4	1850100.000	1978500.000	2058100.000	2133300.000	2257700.000	2369900.000	

	2020	CityName	StateName
0	1713600.000	New York	New York
1	1498400.000	Los Angeles	California
2	443000.000	Houston	Texas
3	893500.000	Chicago	Illinois
4	413700.000	San Antonio	Texas

After evaluation the dataframe, the City and State column are displayed at the last two columns. These need to be brought forward to the beginning.

```
[22]: # Removes Column from data
CityName = result.pop('CityName')
# Inserts Column in first position
result.insert(0, 'CityName', CityName)
# Removes Column from data
StateName = result.pop('StateName')
# Inserts Column in second position
result.insert(1, 'StateName', StateName)
result.head()
```

```
[22]:
```

	CityName	StateName	2008	2009	2010	2011	\
0	New York	New York	0.000	0.000	457300.000	5597200.000	
1	Los Angeles	California	4333600.000	4093300.000	4244700.000	4079000.000	
2	Houston	Texas	1339800.000	1602700.000	1615800.000	1615800.000	
3	Chicago	Illinois	2782900.000	2741900.000	2462400.000	2192000.000	
4	San Antonio	Texas	1309900.000	1588500.000	1576700.000	1579600.000	

	2012	2013	2014	2015	2016	2017	\
0	5697000.000	5804300.000	5870900.000	6219500.000	6444300.000	6510700.000	
1	4165700.000	5326100.000	5515300.000	5917500.000	6367300.000	6817700.000	
2	1679400.000	1897400.000	2025300.000	2051800.000	2171900.000	2293200.000	
3	2178800.000	2536200.000	2867400.000	2840100.000	2631500.000	2722900.000	
4	1628300.000	1738900.000	1850100.000	1978500.000	2058100.000	2133300.000	

	2018	2019	2020
0	6665700.000	6793500.000	1713600.000
1	7667800.000	8478200.000	1498400.000
2	2367300.000	2507100.000	443000.000
3	2960600.000	3234000.000	893500.000
4	2257700.000	2369900.000	413700.000

Now that the years are extracted, the data can be grouped by state and city. This way, the analysis and regression can be done by those two variables. A total can be column will be created as well that will contain the data from all years. The data also has partial months for 2008 and 2020, these two years will be dropped from percentage change so that data with a full years worth will be used and analyzed.

```
[23]: # Groups by State and Sums values
sales_by_state_and_year = result.groupby('StateName').sum()
sales_by_state_and_year.head()
```

```
[23]:
```

	2008	2009	2010	2011 \
StateName				
Alabama	23540500.000	30190100.000	35690200.000	36529000.000
Arizona	82255000.000	84684400.000	79430400.000	73322700.000
Arkansas	10166100.000	13352700.000	15423600.000	18001500.000
California	1062767100.000	1151730100.000	1270672000.000	1254149500.000
Colorado	78392800.000	98927500.000	109680500.000	131729600.000

	2012	2013	2014	2015 \
StateName				
Alabama	46068100.000	58877400.000	70141500.000	83774700.000
Arizona	77522600.000	98266200.000	112274100.000	121594400.000
Arkansas	24754700.000	30623200.000	33655900.000	42085100.000
California	1347562200.000	1646128100.000	1877307500.000	2095765300.000
Colorado	156346900.000	179651600.000	213806800.000	253655300.000

	2016	2017	2018	2019 \
StateName				
Alabama	94137100.000	115069200.000	126179700.000	132996500.000
Arizona	142476000.000	163826400.000	180637800.000	189705100.000
Arkansas	48388200.000	57610900.000	61427300.000	65008500.000
California	2305135300.000	2505317200.000	2815191000.000	2848300300.000
Colorado	295103600.000	339594300.000	377872300.000	394826100.000

	2020
StateName	
Alabama	33213700.000
Arizona	41528300.000
Arkansas	15875900.000

California 608729600.000
 Colorado 98797900.000

```
[24]: # Calculates percentage change on the columns
pct_change_state = sales_by_state_and_year.pct_change(axis=1)*100
pct_change_state.head()
```

```
[24]:      2008    2009    2010    2011    2012    2013    2014    2015    2016  \
StateName
Alabama      NaN  28.247  18.218   2.350  26.114  27.805  19.131  19.437  12.369
Arizona      NaN   2.953  -6.204  -7.689   5.728  26.758  14.255   8.301  17.173
Arkansas      NaN  31.345  15.509  16.714  37.515  23.707   9.903  25.045  14.977
California    NaN   8.371  10.327  -1.300   7.448  22.156  14.044  11.637   9.990
Colorado      NaN  26.195  10.870  20.103  18.688  14.906  19.012  18.638  16.340

      2017    2018    2019    2020
StateName
Alabama    22.236   9.655   5.402  -75.027
Arizona    14.985  10.262   5.020  -78.109
Arkansas    19.060   6.624   5.830  -75.579
California   8.684  12.369   1.176  -78.628
Colorado    15.076  11.272   4.487  -74.977
```

```
[25]: # Calculates total of all sales
sales_by_state_and_year['Total'] = sales_by_state_and_year.sum(axis=1)
sales_by_state_and_year.head()
```

```
[25]:      2008      2009      2010      2011  \
StateName
Alabama    23540500.000  30190100.000  35690200.000  36529000.000
Arizona    82255000.000  84684400.000  79430400.000  73322700.000
Arkansas   10166100.000  13352700.000  15423600.000  18001500.000
California 1062767100.000 1151730100.000 1270672000.000 1254149500.000
Colorado   78392800.000  98927500.000  109680500.000  131729600.000

      2012      2013      2014      2015  \
StateName
Alabama    46068100.000  58877400.000  70141500.000  83774700.000
Arizona    77522600.000  98266200.000  112274100.000  121594400.000
Arkansas    24754700.000  30623200.000  33655900.000  42085100.000
California 1347562200.000 1646128100.000 1877307500.000 2095765300.000
Colorado   156346900.000  179651600.000  213806800.000  253655300.000

      2016      2017      2018      2019  \
StateName
Alabama    94137100.000  115069200.000  126179700.000  132996500.000
Arizona    142476000.000  163826400.000  180637800.000  189705100.000
```

Arkansas	48388200.000	57610900.000	61427300.000	65008500.000
California	2305135300.000	2505317200.000	2815191000.000	2848300300.000
Colorado	295103600.000	339594300.000	377872300.000	394826100.000

	2020	Total
StateName		
Alabama	33213700.000	886407700.000
Arizona	41528300.000	1447523400.000
Arkansas	15875900.000	436373600.000
California	608729600.000	22788755200.000
Colorado	98797900.000	2728385200.000

```
[26]: # Drops the years 2008 and 2020
pct_change_state = pct_change_state.drop(columns=['2008', '2020'])
pct_change_state.head()
```

```
[26]:
```

	2009	2010	2011	2012	2013	2014	2015	2016	2017	\
StateName										
Alabama	28.247	18.218	2.350	26.114	27.805	19.131	19.437	12.369	22.236	
Arizona	2.953	-6.204	-7.689	5.728	26.758	14.255	8.301	17.173	14.985	
Arkansas	31.345	15.509	16.714	37.515	23.707	9.903	25.045	14.977	19.060	
California	8.371	10.327	-1.300	7.448	22.156	14.044	11.637	9.990	8.684	
Colorado	26.195	10.870	20.103	18.688	14.906	19.012	18.638	16.340	15.076	

	2018	2019
StateName		
Alabama	9.655	5.402
Arizona	10.262	5.020
Arkansas	6.624	5.830
California	12.369	1.176
Colorado	11.272	4.487

```
[27]: sales_by_city_and_year = result.groupby('CityName').sum()
sales_by_city_and_year.head()
```

```
[27]:
```

	2008	2009	2010	2011	2012	\
CityName						
Aberdeen	0.000	0.000	0.000	0.000	0.000	
Abilene	0.000	0.000	0.000	0.000	0.000	
Abingdon	2418300.000	2701400.000	2746500.000	2565200.000	2480800.000	
Abington	0.000	0.000	2435100.000	2909200.000	3042500.000	
Accokeek	0.000	0.000	0.000	0.000	0.000	

	2013	2014	2015	2016	2017	\
CityName						
Aberdeen	1109300.000	1922600.000	1925500.000	2036700.000	2119400.000	
Abilene	974300.000	1550500.000	1630700.000	1659100.000	1697300.000	

Abingdon	3050600.000	4473800.000	4564500.000	4485100.000	4568300.000
Abington	3043200.000	3373600.000	3464000.000	3610200.000	3823300.000
Accokeek	0.000	551200.000	3669200.000	3685700.000	3973400.000

	2018	2019	2020
CityName			
Aberdeen	2195300.000	2405500.000	717900.000
Abilene	1754500.000	1940500.000	329400.000
Abingdon	4809700.000	4869600.000	1062000.000
Abington	4077400.000	4186800.000	1083200.000
Accokeek	4220300.000	4340700.000	1112500.000

```
[28]: # Groups by city and adds values
pct_change_city = sales_by_city_and_year.pct_change(axis=1)*100
pct_change_city.head()
```

```
[28]:
```

	2008	2009	2010	2011	2012	2013	2014	2015	2016	2017	\
CityName											
Aberdeen	NaN	NaN	NaN	NaN	NaN	inf	73.317	0.151	5.775	4.060	
Abilene	NaN	NaN	NaN	NaN	NaN	inf	59.140	5.173	1.742	2.302	
Abingdon	NaN	11.707	1.670	-6.601	-3.290	22.968	46.653	2.027	-1.740	1.855	
Abington	NaN	NaN	inf	19.469	4.582	0.023	10.857	2.680	4.221	5.903	
Accokeek	NaN	NaN	NaN	NaN	NaN	NaN	inf	565.675	0.450	7.806	

	2018	2019	2020
CityName			
Aberdeen	3.581	9.575	-70.156
Abilene	3.370	10.601	-83.025
Abingdon	5.284	1.245	-78.191
Abington	6.646	2.683	-74.128
Accokeek	6.214	2.853	-74.370

```
[29]: # Drops years 2008 and 2020
pct_change_city = pct_change_city.drop(columns=['2008', '2020'])
pct_change_city.head()
```

```
[29]:
```

	2009	2010	2011	2012	2013	2014	2015	2016	2017	2018	\
CityName											
Aberdeen	NaN	NaN	NaN	NaN	inf	73.317	0.151	5.775	4.060	3.581	
Abilene	NaN	NaN	NaN	NaN	inf	59.140	5.173	1.742	2.302	3.370	
Abingdon	11.707	1.670	-6.601	-3.290	22.968	46.653	2.027	-1.740	1.855	5.284	
Abington	NaN	inf	19.469	4.582	0.023	10.857	2.680	4.221	5.903	6.646	
Accokeek	NaN	NaN	NaN	NaN	NaN	inf	565.675	0.450	7.806	6.214	

	2019
CityName	
Aberdeen	9.575

```
Abilene 10.601
Abingdon 1.245
Abington 2.683
Accokeek 2.853
```

```
[30]: # Calculates total of all sales
sales_by_city_and_year['Total'] = sales_by_city_and_year.sum(axis=1)
sales_by_city_and_year.head()
```

```
[30]:
```

	2008	2009	2010	2011	2012 \
CityName					
Aberdeen	0.000	0.000	0.000	0.000	0.000
Abilene	0.000	0.000	0.000	0.000	0.000
Abingdon	2418300.000	2701400.000	2746500.000	2565200.000	2480800.000
Abington	0.000	0.000	2435100.000	2909200.000	3042500.000
Accokeek	0.000	0.000	0.000	0.000	0.000

	2013	2014	2015	2016	2017 \
CityName					
Aberdeen	1109300.000	1922600.000	1925500.000	2036700.000	2119400.000
Abilene	974300.000	1550500.000	1630700.000	1659100.000	1697300.000
Abingdon	3050600.000	4473800.000	4564500.000	4485100.000	4568300.000
Abington	3043200.000	3373600.000	3464000.000	3610200.000	3823300.000
Accokeek	0.000	551200.000	3669200.000	3685700.000	3973400.000

	2018	2019	2020	Total
CityName				
Aberdeen	2195300.000	2405500.000	717900.000	14432200.000
Abilene	1754500.000	1940500.000	329400.000	11536300.000
Abingdon	4809700.000	4869600.000	1062000.000	44795800.000
Abington	4077400.000	4186800.000	1083200.000	35048500.000
Accokeek	4220300.000	4340700.000	1112500.000	21553000.000

The final data sets created sales by year for cities and states. This aggregated data can be used to create models for the sales of the city feature and state feature.

Milestone 3

For this project, a linear regression model will be created on the data. This method is being used because it can allow us to uncover patterns and relationships of the data. By exploring the different states, we can see how the predictions trend over time and what assumptions can be made of future data. Grouping the data by year will create different time stamps that can be used in the regression, as well as providing a full year's worth of data to analyze.

```
[31]: # Groups by State and Sums values
StateDf = result.groupby(['StateName'], as_index=False).sum()
StateDf.head()
```

```
[31]:
```

	StateName	2008	2009	2010	2011 \
0	Alabama	23540500.000	30190100.000	35690200.000	36529000.000
1	Arizona	82255000.000	84684400.000	79430400.000	73322700.000
2	Arkansas	10166100.000	13352700.000	15423600.000	18001500.000
3	California	1062767100.000	1151730100.000	1270672000.000	1254149500.000
4	Colorado	78392800.000	98927500.000	109680500.000	131729600.000

		2012	2013	2014	2015	2016 \
0	46068100.000	58877400.000	70141500.000	83774700.000	94137100.000	
1	77522600.000	98266200.000	112274100.000	121594400.000	142476000.000	
2	24754700.000	30623200.000	33655900.000	42085100.000	48388200.000	
3	1347562200.000	1646128100.000	1877307500.000	2095765300.000	2305135300.000	
4	156346900.000	179651600.000	213806800.000	253655300.000	295103600.000	

		2017	2018	2019	2020
0	115069200.000	126179700.000	132996500.000	33213700.000	
1	163826400.000	180637800.000	189705100.000	41528300.000	
2	57610900.000	61427300.000	65008500.000	15875900.000	
3	2505317200.000	2815191000.000	2848300300.000	608729600.000	
4	339594300.000	377872300.000	394826100.000	98797900.000	

```
[32]: # Reshapes data and includes years 2009-2019 to have only full years for the
      ↪ model
years = [str(year) for year in range(2009, 2020)]
StateDf = pd.melt(StateDf, id_vars=['StateName'], value_vars=years,
      ↪ var_name='Year', value_name='Sales')
StateDf.head()
```

```
[32]:
```

	StateName	Year	Sales
0	Alabama	2009	30190100.000
1	Arizona	2009	84684400.000
2	Arkansas	2009	13352700.000
3	California	2009	1151730100.000
4	Colorado	2009	98927500.000

```
[33]: # Drops any NaN to avoid errors
StateDf = StateDf.dropna(axis= 0, how='any')
```

```
[34]: # Separate the target from the features
x = StateDf.drop('Sales', axis=1)
y = StateDf['Sales']
```

```
[35]: from sklearn.linear_model import LinearRegression
```

```
[36]: regression = LinearRegression()
```



```
[37]: # Convert the categorical variables (StateName) into numerical values
X_encoded = pd.get_dummies(x, drop_first=True)
```

```
[39]: # Creates model
model = regression.fit(X_encoded, y)
```

```
[40]: # Creates Predictions
y_pred_test = regression.predict(X_encoded)
```

```
[41]: # Creates dataframe with actual values and predicted values
diff = pd.DataFrame({'Actual': y, 'Predicted': y_pred_test})
```

```
[42]: diff.head()
```

```
[42]:
```

	Actual	Predicted
0	30190100.000	-9918894.620
1	84684400.000	34998069.017
2	13352700.000	-48039067.347
3	1151730100.000	1834408832.653
4	98927500.000	146584832.653

With the model fit and predictions created, the data can be joined to the original names to see how the values differ from the actual. With this dataframe, we can create plots of each state to see the predicted values and actual values graphed to understand how the predictions did.

```
[82]: # Creates dataframe with original years and names
s = pd.concat([x, diff], axis=1)
s.head()
```

```
[82]:
```

	StateName	Year	Actual	Predicted
0	Alabama	2009	30190100.000	-9918894.620
1	Arizona	2009	84684400.000	34998069.017
2	Arkansas	2009	13352700.000	-48039067.347
3	California	2009	1151730100.000	1834408832.653
4	Colorado	2009	98927500.000	146584832.653

```
[98]: # Reshapes the data to have in the original format
ss = s.pivot_table(index='StateName', columns='Year', values=['Actual',
↪ 'Predicted'])
ss.head()
```

```
[98]:
```

		Actual			
	Year	2009	2010	2011	2012
StateName					
Alabama		30190100.000	35690200.000	36529000.000	46068100.000
Arizona		84684400.000	79430400.000	73322700.000	77522600.000
Arkansas		13352700.000	15423600.000	18001500.000	24754700.000

California	1151730100.000	1270672000.000	1254149500.000	1347562200.000
Colorado	98927500.000	109680500.000	131729600.000	156346900.000

Year	2013	2014	2015	2016
StateName				
Alabama	58877400.000	70141500.000	83774700.000	94137100.000
Arizona	98266200.000	112274100.000	121594400.000	142476000.000
Arkansas	30623200.000	33655900.000	42085100.000	48388200.000
California	1646128100.000	1877307500.000	2095765300.000	2305135300.000
Colorado	179651600.000	213806800.000	253655300.000	295103600.000

Year	2017	2018	...	Predicted	
StateName			...	2010	2011
Alabama	115069200.000	126179700.000	...	-2145568.089	3058982.931
Arizona	163826400.000	180637800.000	...	42771395.547	47975946.568
Arkansas	57610900.000	61427300.000	...	-40265740.816	-35061189.796
California	2505317200.000	2815191000.000	...	1842182159.184	1847386710.204
Colorado	339594300.000	377872300.000	...	154358159.184	159562710.204

Year	2012	2013	2014	2015
StateName				
Alabama	17768413.544	45015768.646	67821648.237	92367619.666
Arizona	62685377.180	89932732.282	112738611.874	137284583.302
Arkansas	-20351759.184	6895595.918	29701475.510	54247446.939
California	1862096140.816	1889343495.918	1912149375.510	1936695346.939
Colorado	174272140.816	201519495.918	224325375.510	248871346.939

Year	2016	2017	2018	2019
StateName				
Alabama	117146693.135	144804456.401	171443542.115	182290838.033
Arizona	162063656.772	189721420.037	216360505.751	227207801.670
Arkansas	79026520.408	106684283.673	133323369.388	144170665.306
California	1961474420.408	1989132183.673	2015771269.388	2026618565.306
Colorado	273650420.408	301308183.673	327947269.388	338794565.306

[5 rows x 22 columns]

```
[123]: # Process to create all the plots.

# Calculate the number of rows and columns needed for each plot
num_states = len(ss.index)
num_cols = 6
num_rows = (num_states + num_cols - 1) // num_cols
```

```

# Create a figure and axes for the plots
fig, axes = plt.subplots(nrows=num_rows, ncols=num_cols, figsize=(30,
↳5*num_rows))

# Iterate over the states
for i, state in enumerate(ss.index):
    # Calculate the row and column index for the current state
    row = i // num_cols
    col = i % num_cols

    # Get the data for the current state
    data = ss.loc[state]

    # Create a plot for the current state
    ax = axes[row, col]
    ax.set_title(state)
    ax.set_xlabel('Year')
    ax.set_ylabel('Sales')

    # Plot the Actual values
    data['Actual'].plot(ax=ax, label='Actual')

    # Plot the Predicted values on top of the Actual values
    data['Predicted'].plot(ax=ax, label='Predicted')

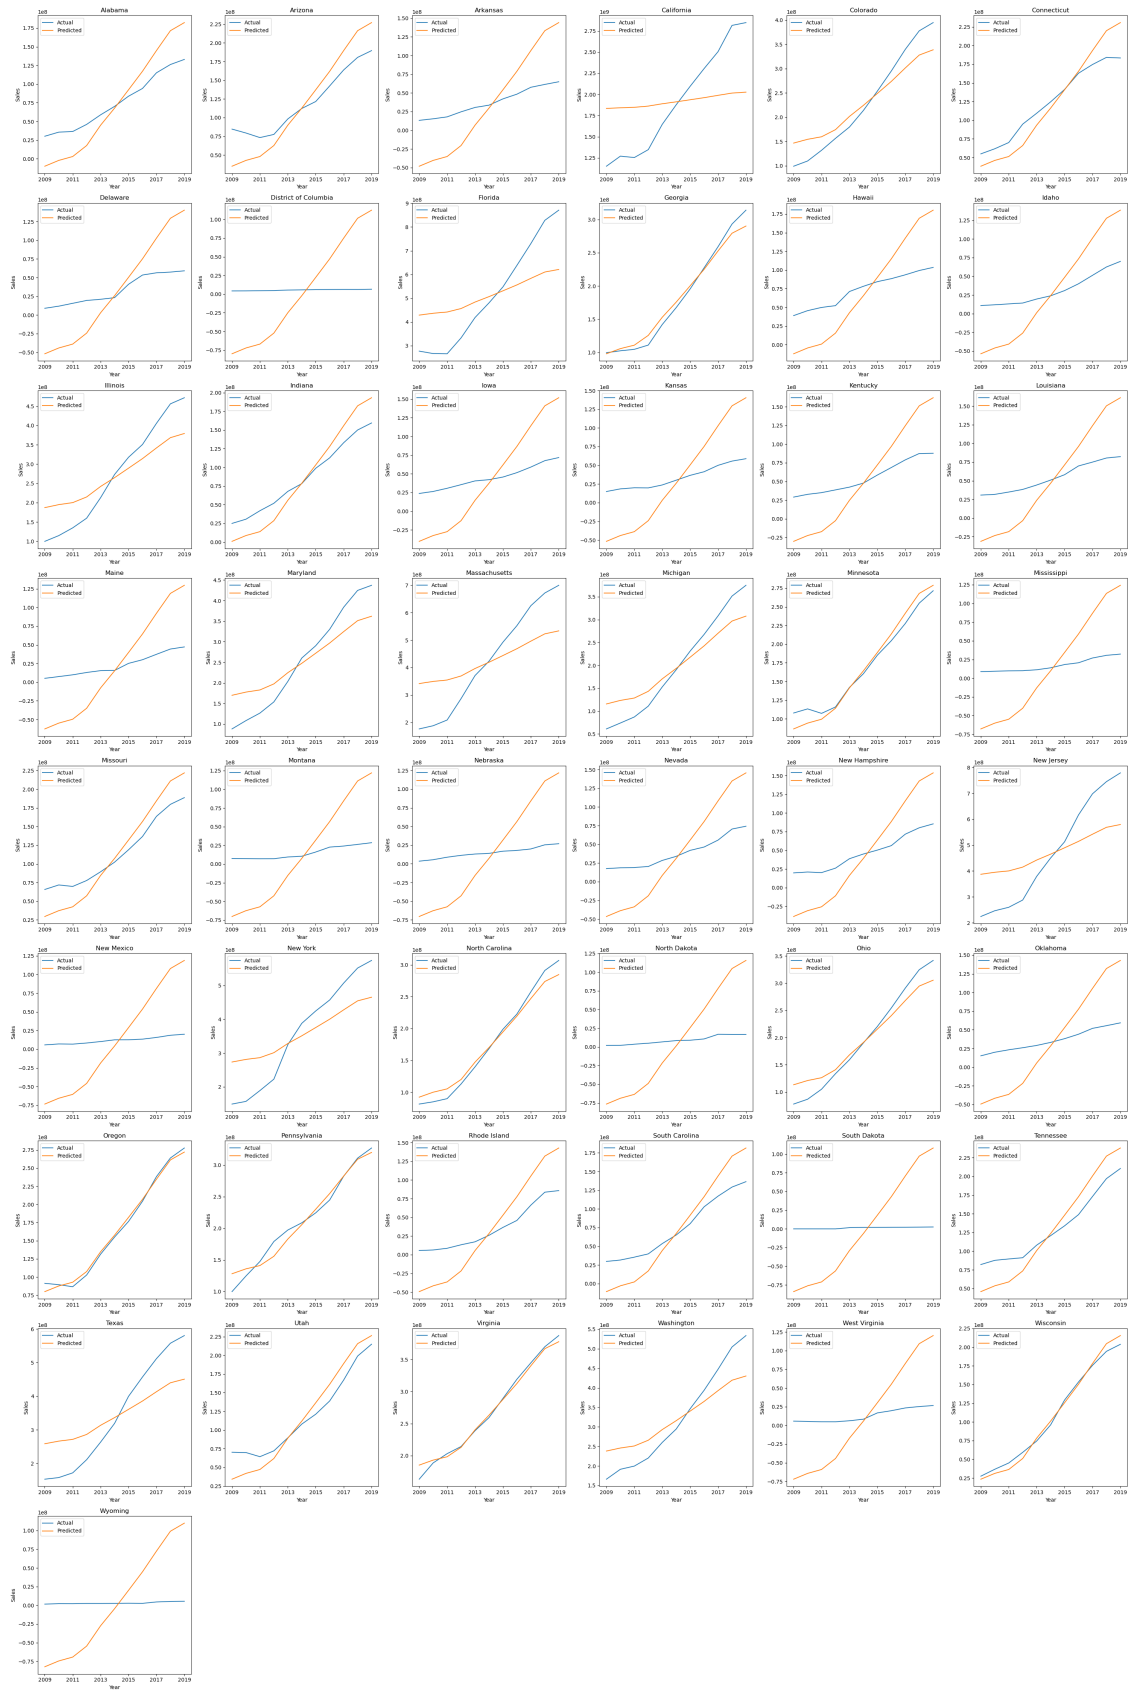
    # Add a legend
    ax.legend()

# Remove any unused axes
if num_states < num_rows * num_cols:
    for i in range(num_states, num_rows * num_cols):
        fig.delaxes(axes.flatten()[i])

# Adjust the layout and spacing
plt.tight_layout()

# Show the plots
plt.show()

```



The plots above show the actual numbers and predictions by state. At a glance, it can be seen that many of the states predicted values are similar to the actual values. Also, almost all fit have the same positive trend where they increase sales over time.

```
[108]: from sklearn import metrics
      from sklearn.metrics import r2_score
```

```
[107]: r2 = r2_score(y, y_pred_test)
      mae = metrics.mean_absolute_error(y, y_pred_test)
```

```
[106]: print('Accuracy:', r2.round(2)*100, '%')
      print('MAE:', mae.round(2))
```

Accuracy: 90.0 %

MAE: 51787151.26

The accuracy of the model was 90%. We can see from this model that many of the states follow a positive trend over time. Due to the high accuracy, we can see the predictions and actual sales are similar in most states. This model was able to accurately represent some of the states but with the data, there was some error to be expected. The MAE value tells us the average error is around 51 million, but when looking at sales, many lie in the 100s of millions or more, so this is a good value.

Data Source: <https://www.kaggle.com/datasets/paultimothymooney/zillow-house-price-data/discussion>