

Week 1

Felipe Rodriguez

Bellevue University

DSC 650 Big Data

Professor Nasheb Ismaily

December 3, 2023

The screenshot shows a web browser window with the address bar set to "localhost". The page title is "Overview" for "Hadoop 3007" (active). The navigation bar includes links like Home, Overview, Datanodes, Datanode Volume Failures, Snapshot, Backup Progress, Uptime, etc.

Overview

Status:	Sun Dec 03 12:34:02 -0800 2023
Version:	3.2.3-noble56f1d37c0095488613d598bc8e5f4d9f31
Completed:	Sat Mar 19 20:18:02 -0800 2022 by ubuntu from branch-3.2.3
Cluster ID:	CID-48652639-6143-4040-a9dc-d5ac98198709
Block Pool ID:	BP-203c703294-172.28.1.1-1701628398223

Summary

Security is off.
Safemode is off.
1 file and directories, 0 blocks (0 replicated blocks, 0 erasure coded block groups) = 1 total (filesystem objects)
Heap Memory used 118.14 MB of 241.5 MB Heap Memory. Max Heap Memory is 1.72 GB.
Non-Heap Memory used 10.27 MB of 51.56 MB Committed Non-Heap Memory. Max Non-Heap Memory is configured.
Configured Capacity: 96 GB
Configured Remote Capacity: 0 B
DFS Used: 49 KB (0%)
Non DFS Used: 16.8B GB
DFS Remaining: 79.66 GB (82.5%)
Block Pool Used: 49 KB (0%)
Datanode usage: 0.0% / 16.80% / 0.00%
Live Nodes: 2 (Decommissioned: 0, In Maintenance: 0)
Dead Nodes: 2 (Decommissioned: 0, In Maintenance: 0)
Decommissioning Nodes: 0
Entering Maintenance Nodes: 0
Total Datanode Volume Failures: 0 (0 B)
Number of Under-Replicated Blocks: 0
Number of Blocks Pending Retention (including replicat): 0
Block Deletion Start Time: Sun Dec 03 12:34:02 -0800 2023
Last Checkpoint Time: Sun Dec 03 12:33:19 -0800 2023
Enabled Erasure Coding Policies: RS-0-3-12GB

NameNode Journal Status


Current transaction ID: 1
Journal Manager: State
FileJournalManager(inodeName): EditLogOutputStream(fileName=namenode_1_inprogress_00000000000000000001)

NameNode Storage

Storage Directory	Type	State
/dfs/name	RANGE_AND_EDITS	Active

NFS Name Space

Spark Master



Spark 3.0.0 Spark Master at spark/master:7077


URL: http://localhost:7077
Spark Version: 3.0.0 (3.0.0)
Spark UI: http://localhost:7077
Memory in use: 12.5 GB Total, 0.0 GB Used
Approximate: 100% (0.0%)
Spark UI: http://localhost:7077
Spark UI: http://localhost:7077

Worker ID	Address	State	Cores	Memory	Resources
worker-0000	worker-0000	ALIVE	4 GB Used	6.7 GB (0.0 GB Used)	
worker-0001	worker-0001	ALIVE	4 GB Used	6.7 GB (0.0 GB Used)	

Application ID	Name	Cores	Memory per Executor	Resources Per Executor	Submitted Time	User	State	Duration
----------------	------	-------	---------------------	------------------------	----------------	------	-------	----------

Application ID	Name	Cores	Memory per Executor	Resources Per Executor	Submitted Time	User	State	Duration
----------------	------	-------	---------------------	------------------------	----------------	------	-------	----------

Spark History



Spark 3.0.0 History Server

Event log directory: http://localhost:18080
Last updated: 2020-12-03 12:38:58
Client local time zone: America/Chicago
No completed applications found!
Did you specify the correct logging directory? Please verify your setting of spark.history.logDirectory listed above and whether you have the permissions to access it.
It is also possible that your application did not run to completion or did not stop the SparkContext.
[Show incomplete applications](#)

NIFI

