

Week 4

July 2, 2023

Load the data as a Pandas data frame and ensure that it imported correctly.

```
[233]: import pandas as pd
```

```
[234]: # Read Data
data = pd.read_csv('auto-mpg.csv')
data.head()
```

```
[234]:
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	model year	\
0	18.0	8	307.0	130	3504	12.0	70	
1	15.0	8	350.0	165	3693	11.5	70	
2	18.0	8	318.0	150	3436	11.0	70	
3	16.0	8	304.0	150	3433	12.0	70	
4	17.0	8	302.0	140	3449	10.5	70	

	origin	car name
0	1	chevrolet chevelle malibu
1	1	buick skylark 320
2	1	plymouth satellite
3	1	amc rebel sst
4	1	ford torino

Begin by prepping the data for modeling:

Remove the car name column.

```
[236]: # Look at name of columns
data.columns
```

```
[236]: Index(['mpg', 'cylinders', 'displacement', 'horsepower', 'weight',
        'acceleration', 'model year', 'origin', 'car name'],
        dtype='object')
```

```
[237]: # Remove column
data = data.drop(columns="car name")
```

```
[238]: data.head()
```

```
[238]:      mpg  cylinders  displacement  horsepower  weight  acceleration  model year  \
0   18.0         8         307.0         130    3504         12.0         70
1   15.0         8         350.0         165    3693         11.5         70
2   18.0         8         318.0         150    3436         11.0         70
3   16.0         8         304.0         150    3433         12.0         70
4   17.0         8         302.0         140    3449         10.5         70

      origin
0         1
1         1
2         1
3         1
4         1
```

The horsepower column values likely imported as a string data type. Figure out why and replace any strings with the column mean.

```
[239]: # Display data types
data.dtypes
```

```
[239]: mpg          float64
cylinders        int64
displacement     float64
horsepower       object
weight           int64
acceleration     float64
model year       int64
origin           int64
dtype: object
```

```
[241]: # Changes horsepower to numeric and converts errors to NaN
data['horsepower'] = pd.to_numeric(data['horsepower'], errors='coerce')
```

```
[243]: # Removes any NaN
data = data.dropna(axis= 0, how='any')
```

```
[244]: data.dtypes
```

```
[244]: mpg          float64
cylinders        int64
displacement     float64
horsepower       float64
weight           int64
acceleration     float64
model year       int64
origin           int64
dtype: object
```

Create dummy variables for the origin column.

```
[245]: # Creates dummies for origin column
dummy_origin = pd.get_dummies(data["origin"])
```

```
[246]: dummy_origin
```

```
[246]:
```

	1	2	3
0	1	0	0
1	1	0	0
2	1	0	0
3	1	0	0
4	1	0	0
..
393	1	0	0
394	0	1	0
395	1	0	0
396	1	0	0
397	1	0	0


[392 rows x 3 columns]

Create a correlation coefficient matrix and/or visualization. Are there features highly correlated with mpg?

```
[247]: # Creates correlation matrix
corr_matrix = data.corr()
```

```
[248]: import seaborn as sns
```

```
[249]: # Displays correlation matrix
cm_plot = sns.heatmap(corr_matrix, annot=True)
```



The image above created from the correlation matrix of the Auto Data, shows us that MPG is highly correlated with the following features: Displacement and Weight.


Plot mpg versus weight. Analyze this graph and explain how it relates to the corresponding correlation coefficient.

```
[250]: # Import Necessary Libraries
import matplotlib as plot
```

```
[251]: # Converts weight to numeric to avoid errors
data['weight'] = pd.to_numeric(data['weight'], errors='coerce')
```

```
[252]: # Creates scatter plot of age versus weight
data.plot.scatter(x = 'mpg', y = 'weight')
```

```
[252]: <AxesSubplot:xlabel='mpg', ylabel='weight'>
```



MPG versus weight has a correlation coefficient of - 0.83. Since this is so close to - 1, it indicates that there is a negative correlation between the two. As weight goes down, MPG goes up. The scatter plot shows this as the points are in a negative slope as MPG increases

Randomly split the data into 80% training data and 20% test data, where your target is mpg.

```
[253]: from sklearn.model_selection import train_test_split
```

```
[254]: # Separate the target from the features
feature = data.drop('mpg', axis=1)
target = data['mpg']

#Split the data into 80% training and 20% test
feature_train, feature_test, target_train, target_test = \
    ↪train_test_split(feature, target, test_size=0.2, random_state=42)
```

Train an ordinary linear regression on the training data.

```
[255]: # Import necessary Libraries
from sklearn.linear_model import LinearRegression
```

```
[256]: # Create Linear Regression Model
regression = LinearRegression()
```

```
[257]: # Drops any NaN to avoid errors
data = data.dropna(axis= 0, how='any')
```

```
[258]: # Creates model
model = regression.fit(feature_train, target_train)
```

Calculate R2, RMSE, and MAE on both the training and test sets and interpret your results.

```
[259]: # R2 for train data
regression.score(feature_train, target_train)
```

```
[259]: 0.826001578671067
```

```
[260]: # R2 for test data
# R2 for train data
regression.score(feature_test, target_test)
```

```
[260]: 0.7901500386760352
```

```
[261]: # Import necessary libraries
from sklearn.metrics import mean_squared_error
```

```
[262]: model.fit(feature_test, target_test)
```

```
[262]: LinearRegression()
```

```
[263]: # Creates predictions of test features
y_pred_test = model.predict(feature_test)
```

```
[264]: # RMSE of test
rmse_test = mean_squared_error(target_test, y_pred_test)**0.5
rmse_test
```

```
[264]: 2.9816057261316895
```

```
[265]: model.fit(feature_train, target_train)
```

```
[265]: LinearRegression()
```

```
[266]: # Creates predictions of training features
y_pred_train = model.predict(feature_train)
```

```
[267]: # RSEM of train
rmse_train = mean_squared_error(target_train, y_pred_train)**0.5
```

```
rmse_train
```

```
[267]: 3.3134960151437447
```

```
[268]: # Import necessary Libraries  
from sklearn.metrics import mean_absolute_error as mae
```

```
[269]: # MAE of Test  
print("MAE Test:", mae(target_test,y_pred_test))
```

```
MAE Test: 2.1864311671060714
```

```
[270]: # MAE of Train  
print("MAE Train:", mae(target_train,y_pred_train))
```

```
MAE Train: 2.548168196215135
```

Pick another regression model and repeat the previous two steps. Note: Do NOT choose logistic regression as it is more like a classification model.

```
[271]: # Import necessary Libraries  
from sklearn.linear_model import Ridge
```

```
[272]: # Create Ridge Regression  
ridgeReg = Ridge(alpha=10)  
  
ridgeReg.fit(feature_train,target_train)
```

```
[272]: Ridge(alpha=10)
```

```
[273]: # Runs ridge regression  
train_score_ridge = ridgeReg.score(feature_train, target_train)  
test_score_ridge = ridgeReg.score(feature_test, target_test)
```

```
[274]: print("The train score for ridge model is {}".format(train_score_ridge))  
  
print("The test score for ridge model is {}".format(test_score_ridge))
```

```
The train score for ridge model is 0.8258936780257804
```

```
The test score for ridge model is 0.7915211170468783
```